

Computer Networks

COL 334/672

Software Defined Networking

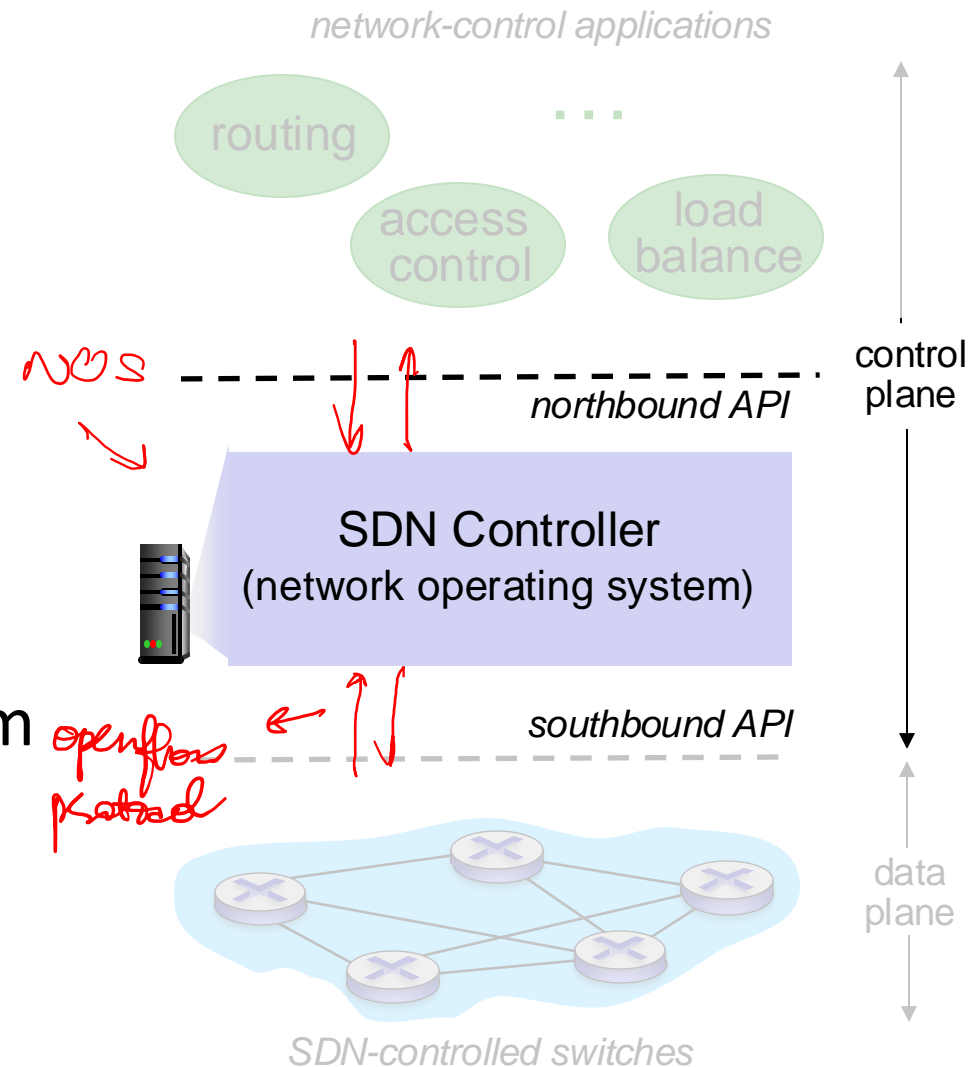
Tarun Mangla

Slides adapted from KR

Sem 1, 2024-25

Recap: SDN Controller

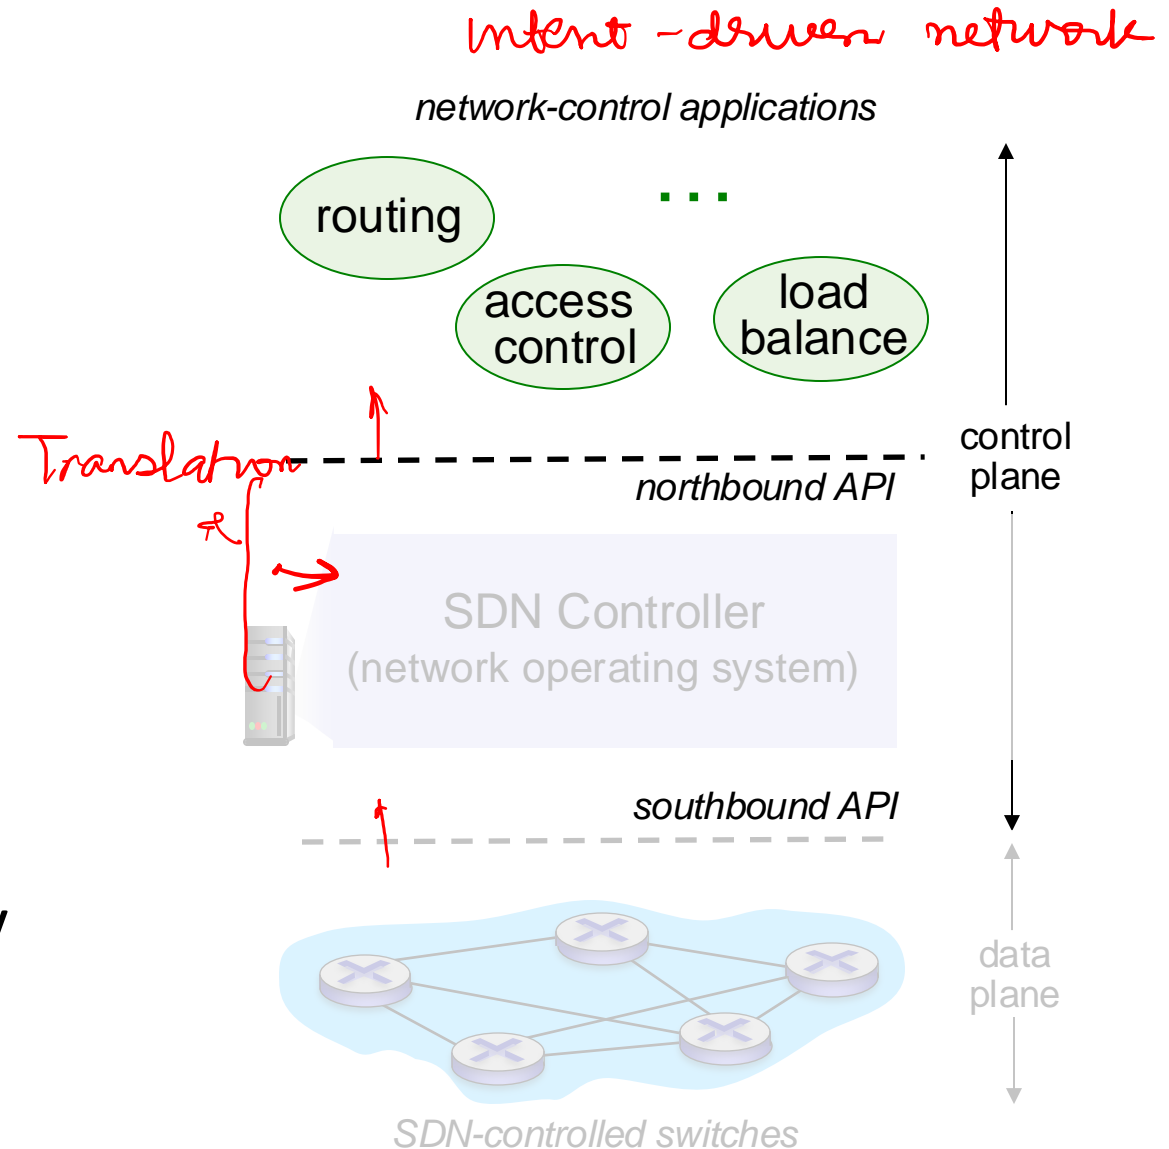
- maintains network state information
- interacts with network control applications “above” via northbound API
- interacts with network switches “below” via southbound API
- implemented as distributed system for performance, scalability, fault-tolerance, robustness



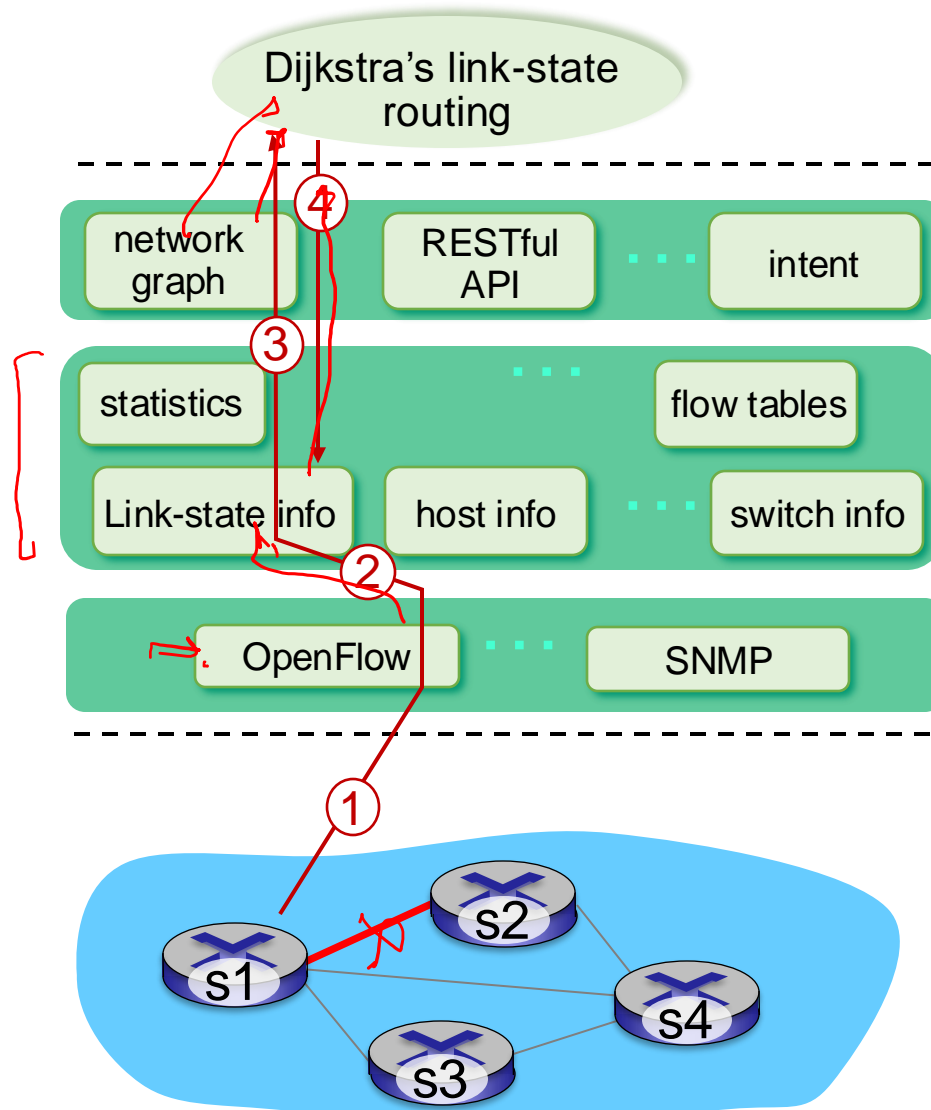
protocol → application

Software defined networking (SDN)

- operators don't "program" switches by creating/sending OpenFlow messages directly.
- Instead use higher-level abstraction at controller
- "brains" of control: implement control functions using lower-level services, API provided by SDN controller
- *unbundled*: can be provided by 3rd party: distinct from routing vendor, or SDN controller

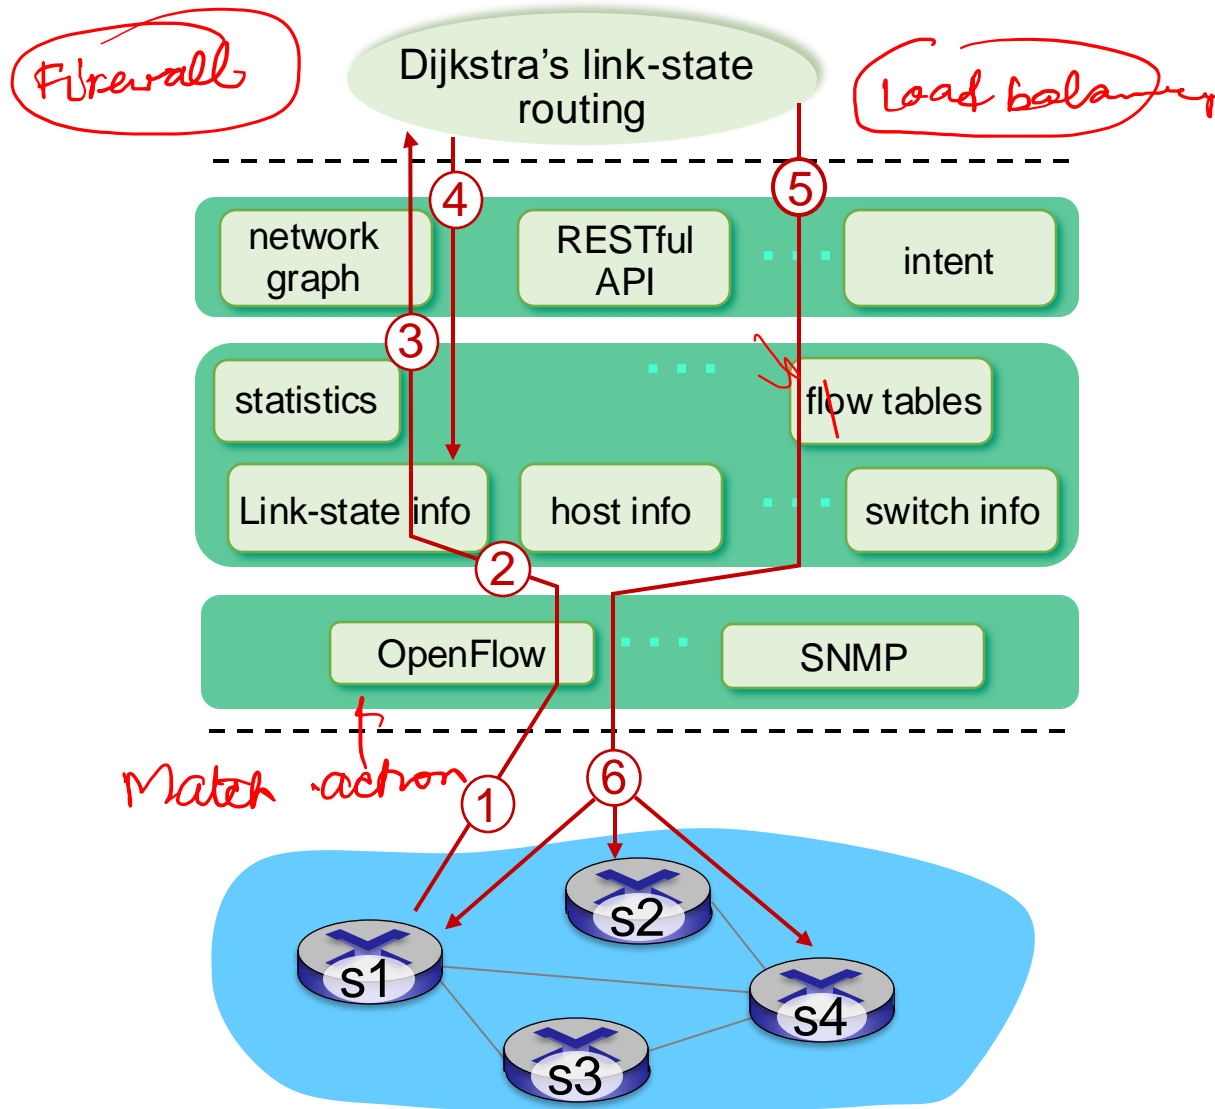


SDN: control/data plane interaction example



- ① S1, experiencing link failure uses OpenFlow port status message to notify controller
- ② SDN controller receives OpenFlow message, updates link status info
- ③ Dijkstra's routing algorithm application has previously registered to be called when ever link status changes. It is called.
- ④ Dijkstra's routing algorithm access network graph info, link state info in controller, computes new routes

SDN: control/data plane interaction example

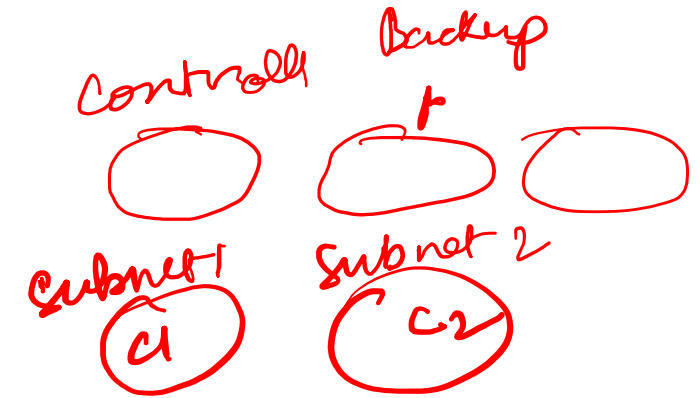


- ⑤ link state routing app interacts with flow-table-computation component in SDN controller, which computes new flow tables needed
- ⑥ controller uses OpenFlow to install new tables in switches that need updating


SDN: Key Challenges

■ Hardening the control plane

- Scalability
 - Reliability
 - Consistency
 - Security
- Real-time → logically centralized
- Distributed systems



■ Internet-scaling: beyond a single AS (?)

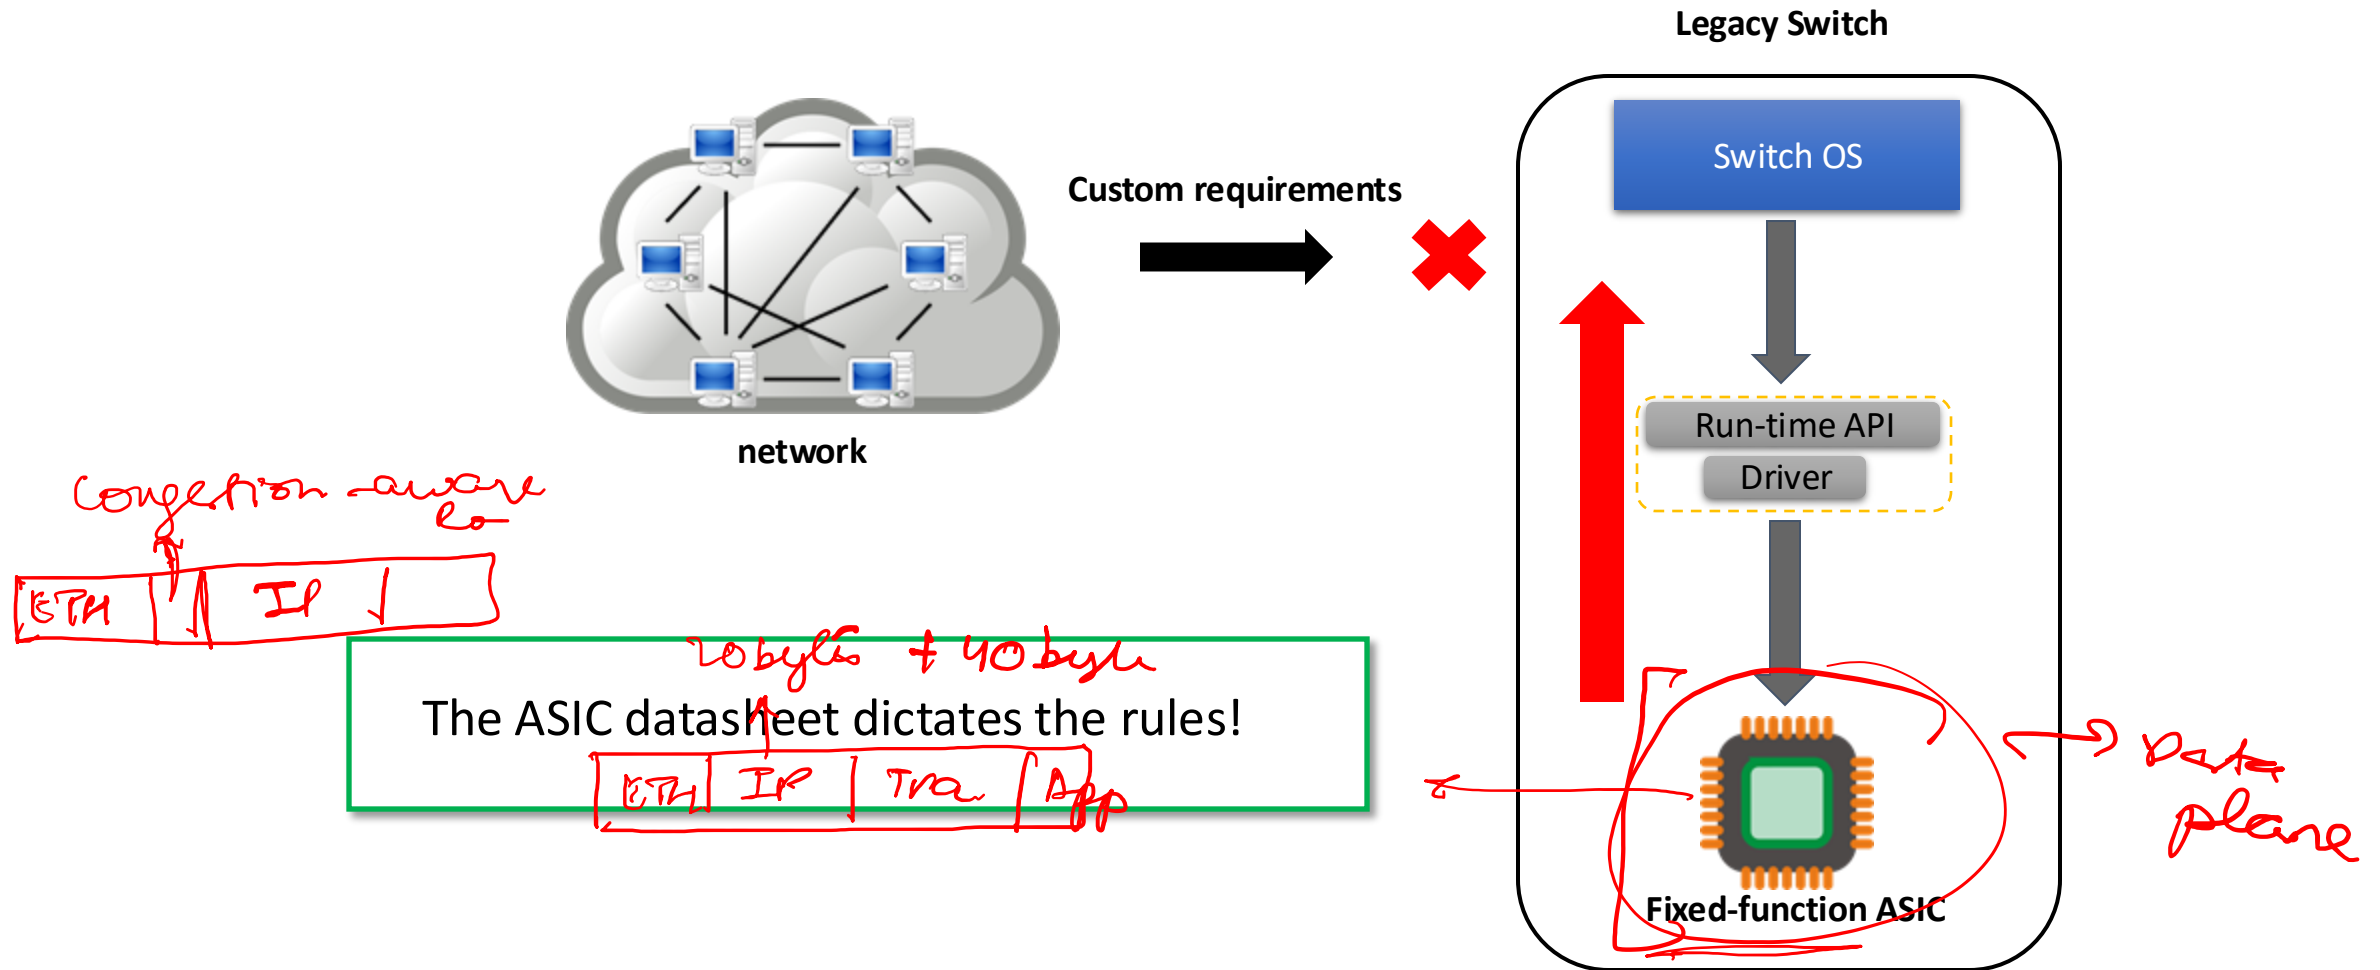



What else is programmable in the network?

- Programmable data plane }
- Network function virtualization (NFV)

↓
making middleboxes programmable

Status Quo: Bottom-up design


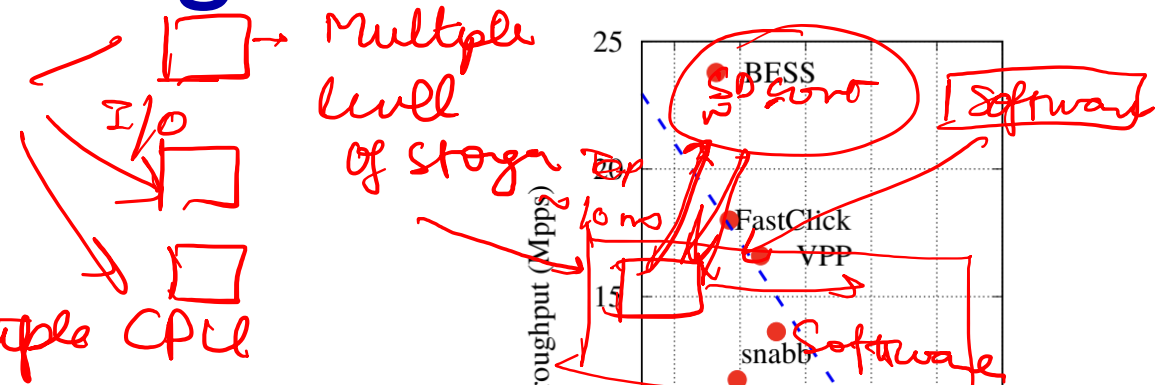





How to make Data Plane Programmable?


■ Move the data plane to software

- Generally, too slow for data plane functions
- Any optimization techniques?
 - Parallelism across multiple servers and cores
 - Optimizations in NUMA
 - Fast I/O ...



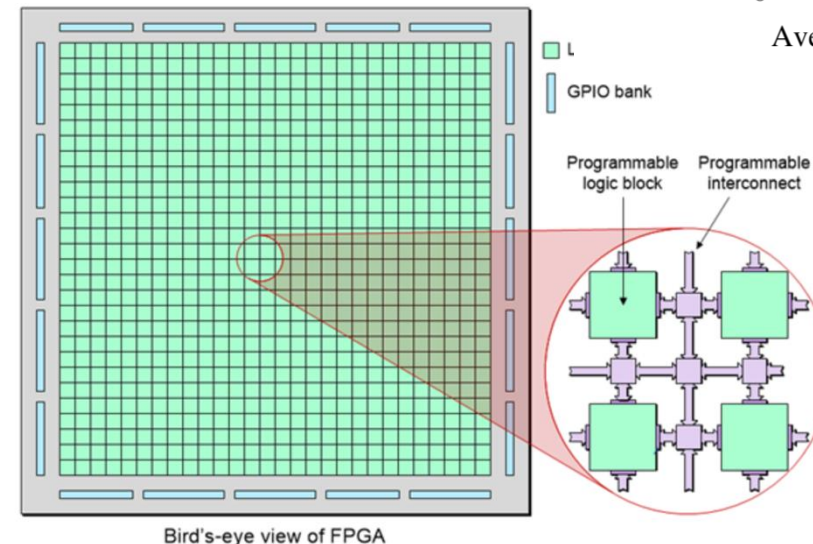
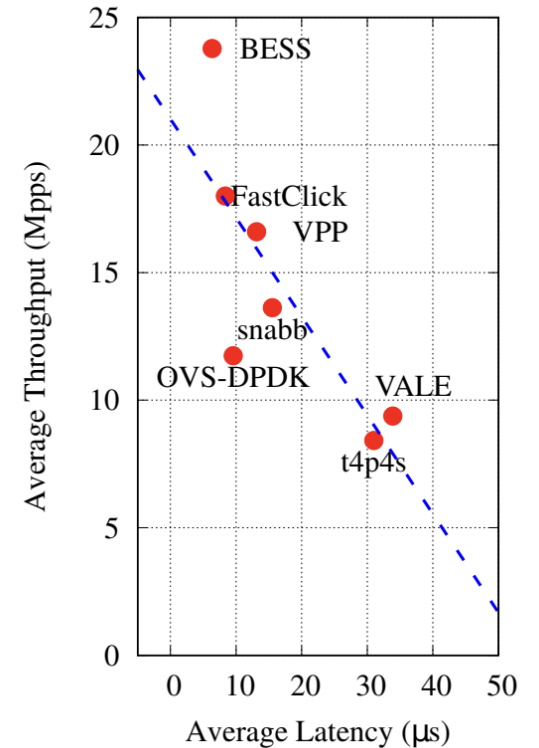
■ What about programmable hardware?

- FPGA: but costly, power-hungry, slower



How to make Data Plane Programmable?

- Move the data plane to software
 - Generally, too slow for data plane functions
 - Any optimization techniques?
 - Parallelism across multiple servers and cores
 - Optimizations in NUMA
 - Fast I/O ...
- What about programmable hardware?
 - FPGA: but costly, power-hungry, slower

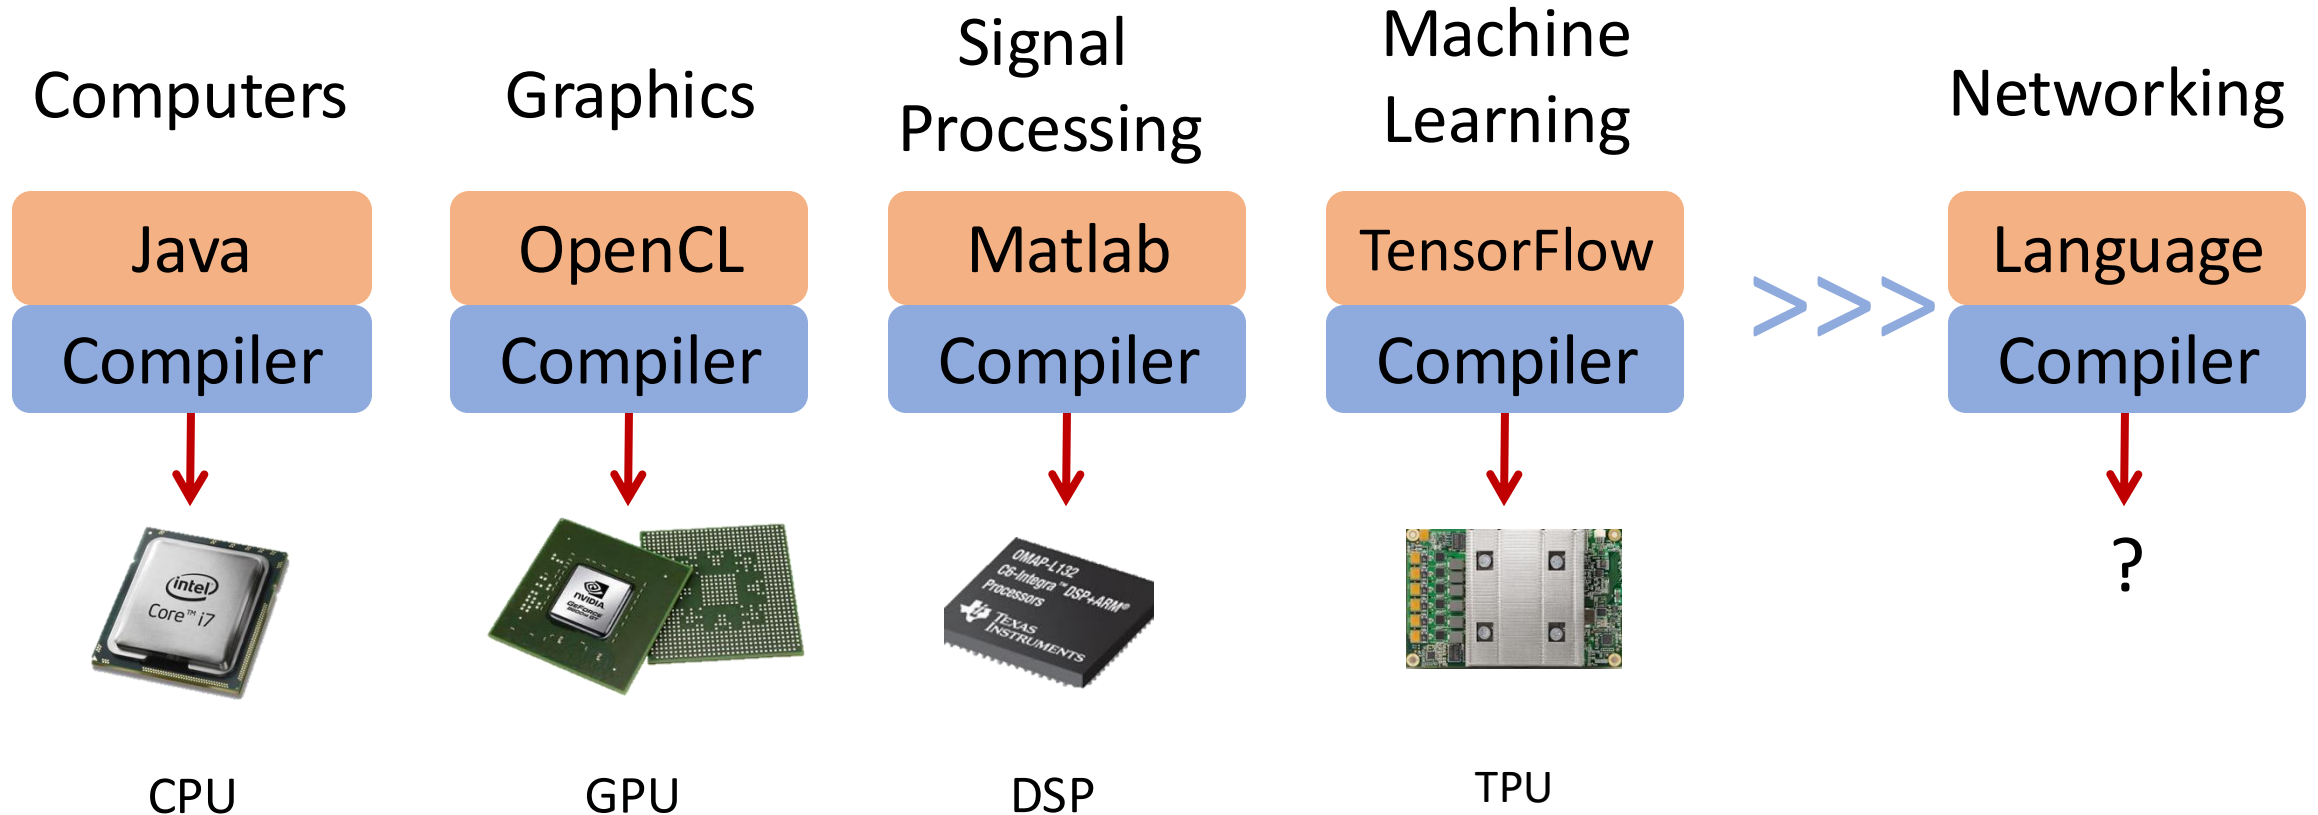


What about Programmable Hardware?

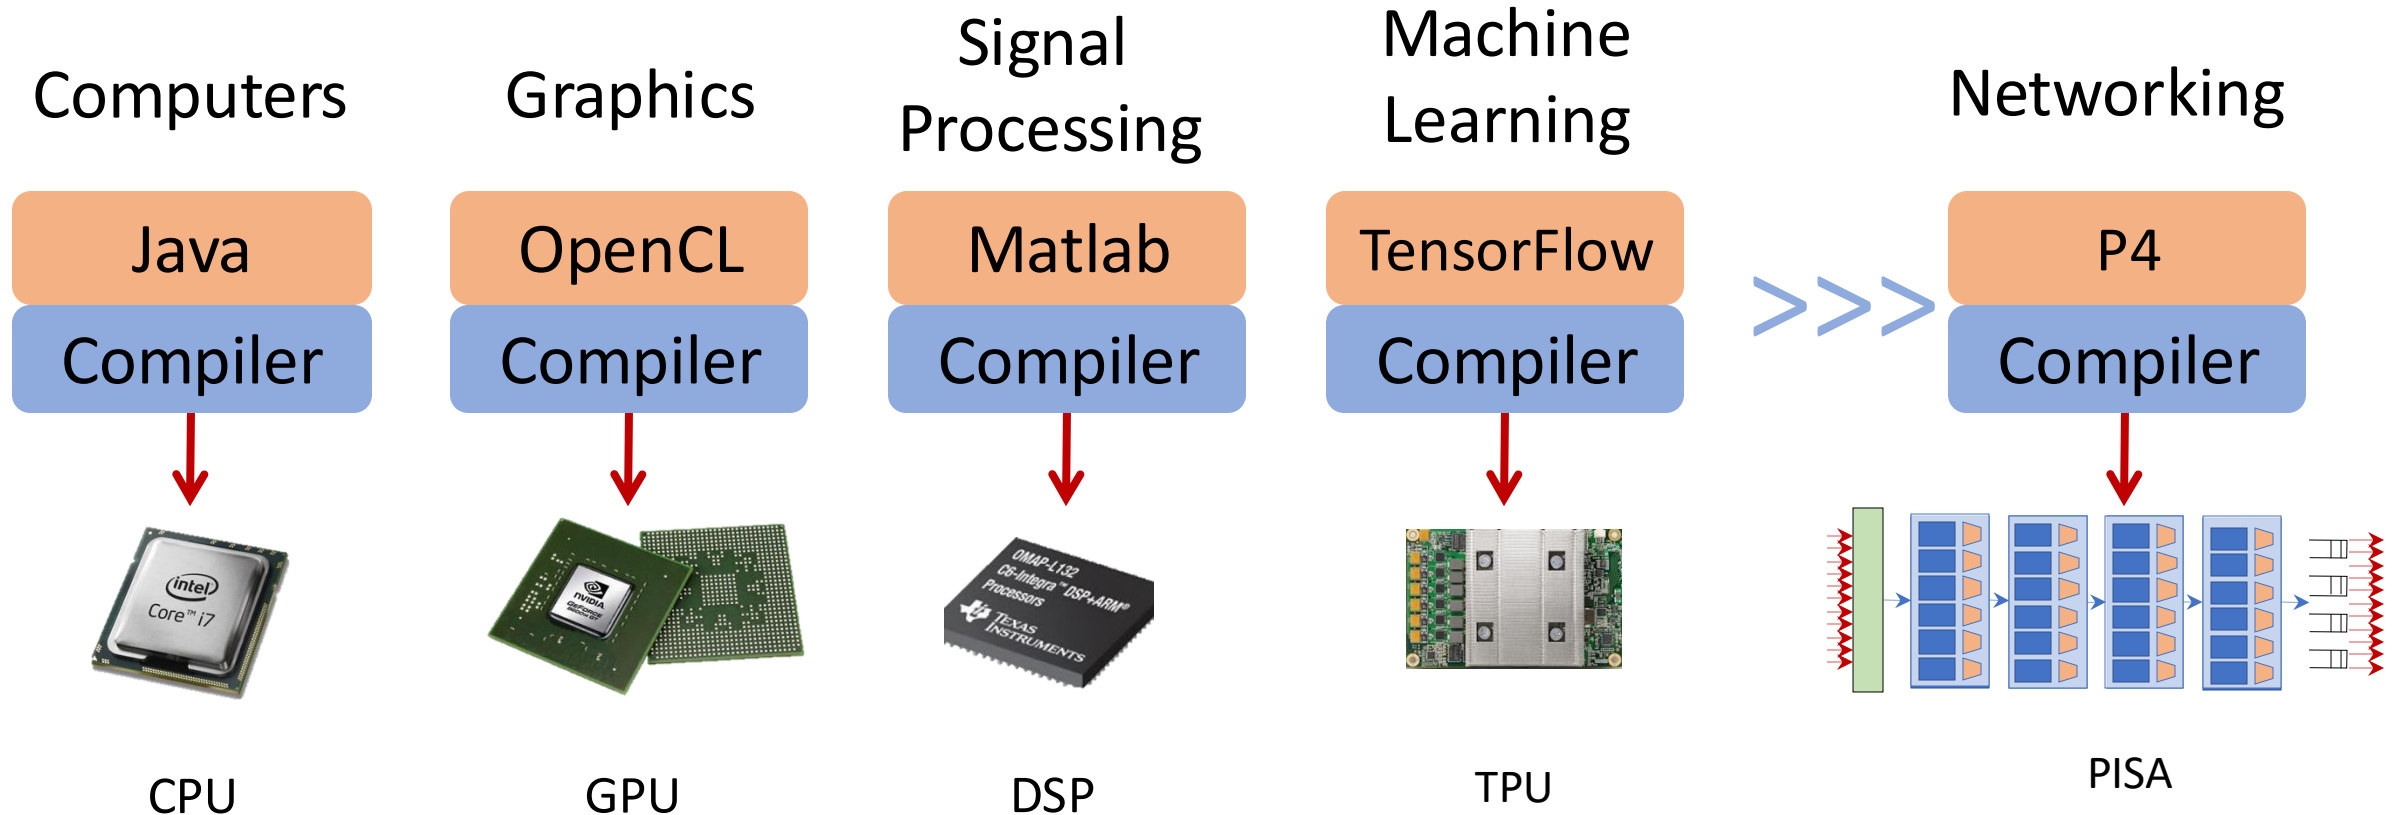
“Programmable switches are 10-100x slower than fixed-function switches. They are more expensive and consume more power.”

Conventional wisdom in networking


Domain Specific Processors



Domain Specific Processors



Protocol Independent Switch Architecture



P4 describes PISA's behavior

Reducing complexity

switch.p4

Switch OS

IPv4 and IPv6 routing

- Unicast Routing
 - Routed Ports & SVI
 - VRF
- Unicast RPF
 - Strict and Loose

~~Multicast~~

- ~~- PIM-SM/DM & PIM-Bidir~~

Ethernet switching

- ~~- VLAN Flooding~~
- MAC Learning & Aging
- STP state
- ~~- VLAN Translation~~

Load balancing

- ~~- LAG~~
- ECMP & WCMP
- Resilient Hashing
- ~~- Flowlet Switching~~

Fast Failover

- LAG & ECMP

Tunneling

- IPv4 and IPv6 Routing & Switching
 - ~~- IP in IP (Gin4, 4in4)~~
 - VXLAN, NVGRE, GENEVE & GRE
 - ~~- Segment Routing, ILA~~

~~MPLS~~

- ~~- LER and LSR~~
- ~~- IPv4/v6 routing (L3VPN)~~
- ~~- L2 switching (EoMPLS, VPLS)~~
- ~~- MPLS over UDP/GRE~~

ACL

- MAC ACL, IPv4/v6 ACL, RACL
- ~~- QoS ACL, System ACL, PBR~~
- Port Range lookups in ACLs

QoS

- QoS Classification & marking
- ~~- Drop profiles/WRED~~
- ~~- RoCE v2 & FCoE~~
- CoPP (Control plane policing)

~~NAT and L4 Load Balancing~~

Security Features

- ~~- Storm Control, IP Source Guard~~

Monitoring & Telemetry

- ~~- Ingress Mirroring and Egress Mirroring~~
- Negative Mirroring
- ~~- Sflow~~
- INT

Counters

- Route Table Entry Counters
- ~~- VLAN/Bridge Domain Counters~~
- Port/Interface Counters

Protocol Offload

- BFD, OAM

Multi-chip Fabric Support

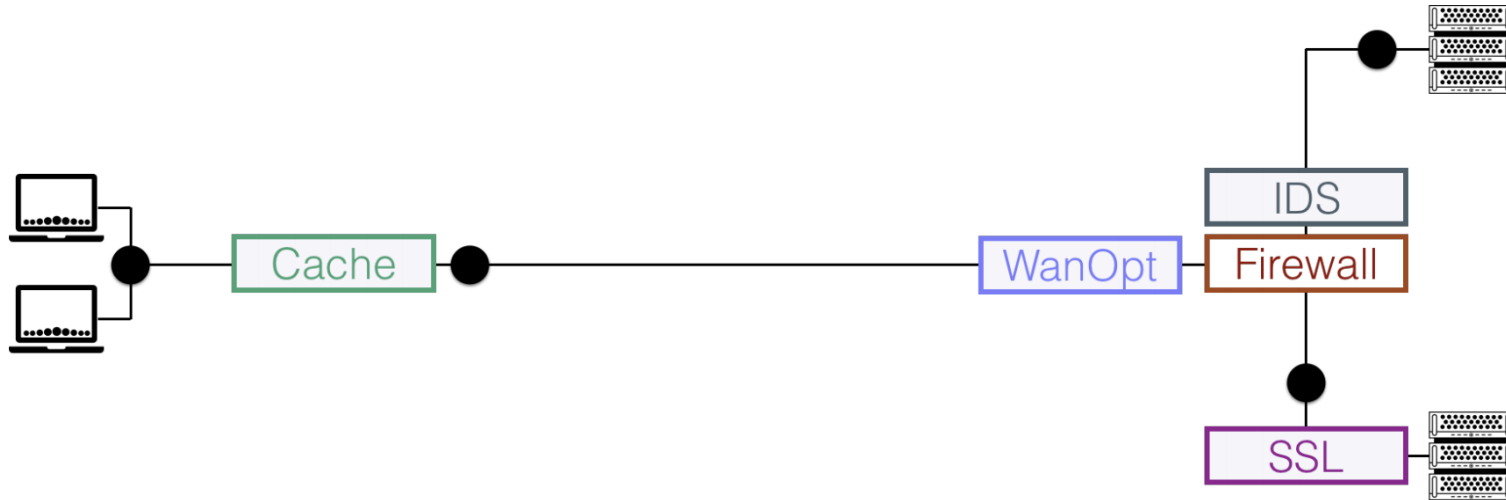
- ~~- Forwarding, QoS~~

What else is programmable in the network?

- Programmable data plane
- **Network function virtualization**

Middleboxes

Data delivery is not the only required functionality.




Elements in the network path for security, performance enhancements etc.


One-third of all network devices in enterprises are middleboxes!
Sherry et al., SIGCOMM'12

Evolution of Middleboxes


Dedicated hardware



*Need for
flexibility*


A black arrow pointing from the dedicated hardware middlebox to the software middlebox.

Software

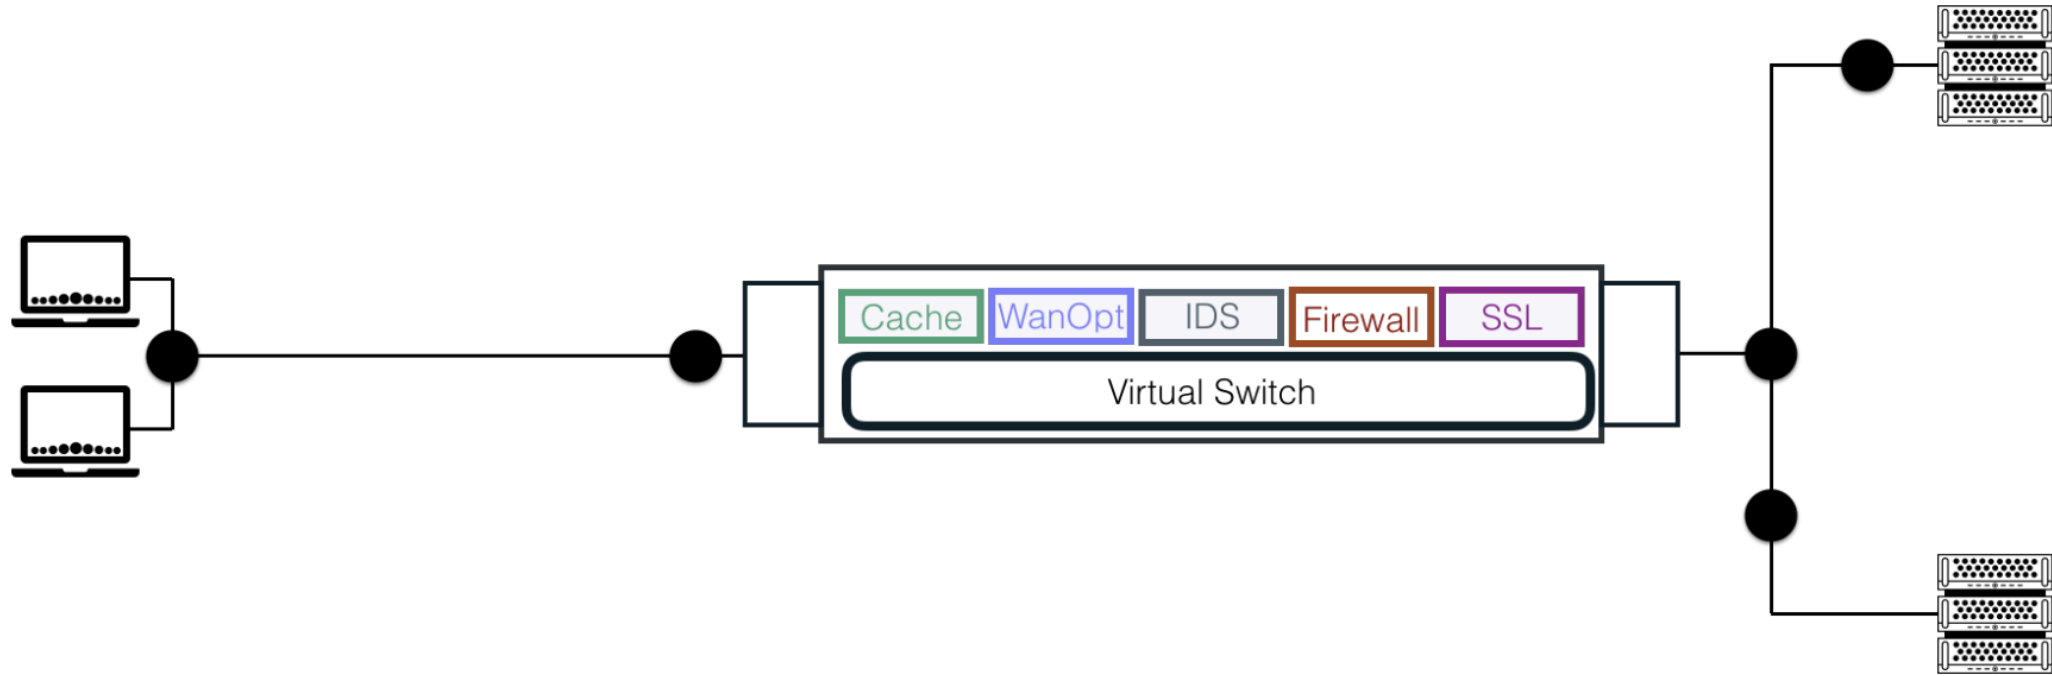


From Hardware Middleboxes..

Data delivery is not the only required functionality.



To Software Network Functions (NF)




Primarily deployed in a VM
(Network Function Virtualization or NFV)

Functional Elements, not Middleboxes

- **WAN Optimizer** = Caching + Deduplication + Compression + Encryption + Forward Error Correction + Rate Limiter
- **Application Firewall** = IP Defragmenter + Application Detection Engine + Logger + Blocker
- **IDS** = IP Defragmenter + Preprocessing + Misuse Detection Engine + Logger

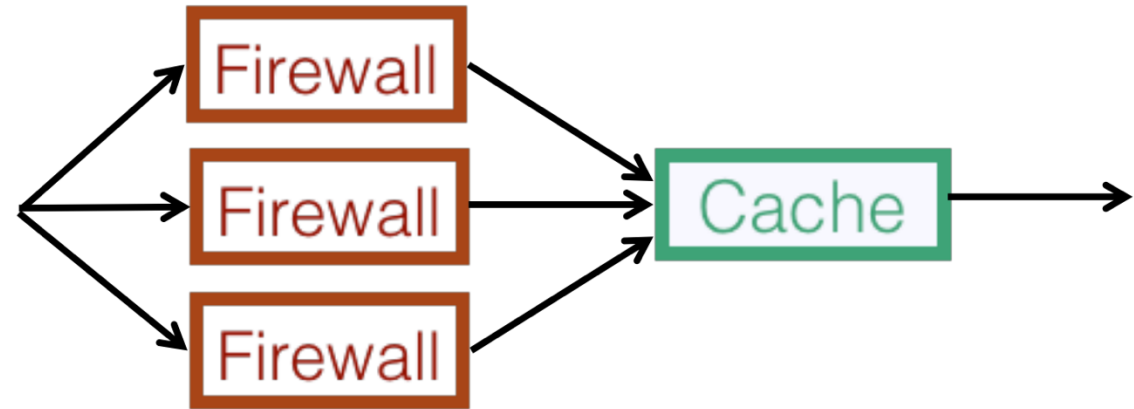
Why NFV?

- Softwarization leads to faster innovation



Why NFV?

- Softwarization leads to faster innovation
- Ease of deployment, configuration, and management
- Consolidation: Reduce number of hardware boxes in the network



Being adopted by both carriers and cloud providers

NFV Challenges

- Virtual network function management
 - Where and how to install network functions?
- Unpredictable (low) Performance
 - How to mitigate the overheads of virtualization?
- Fault Tolerance
 - How to handle recovery in case of faults?

Summary

- Increased programmability in the networks
 - Greater flexibility → Faster innovation
- Programmable control plane
 - SDN
- Programmable data plane
 - SDN-2 / P4
- Software middleboxes implemented in VMs
 - Network function virtualization (NFV)

Attendance

