

Computer Networks

COL 334/672

Transport Layer

Tarun Mangla

Slides adapted from KR

Sem 1, 2024-25

Recap


■ Transport-layer services

- Multiplexing/demultiplexing
- Reliability
- Flow Control
- Congestion control

} UDP
TCP


■ Reliability

- Automatic Repeat reQuest (ARQ) protocols
 - Stop-and-wait protocol
 - Sliding window protocol

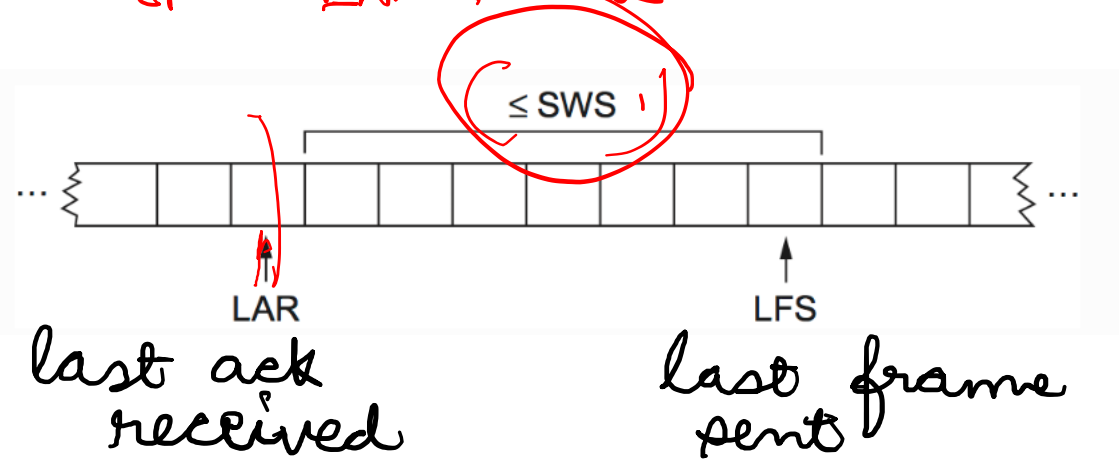


Sliding Window Protocol

Sender window size




$$LFS - LAR \leq SWS$$



RWS


$$RWS = 1$$

$$(LAF - LFR) \leq RWS$$



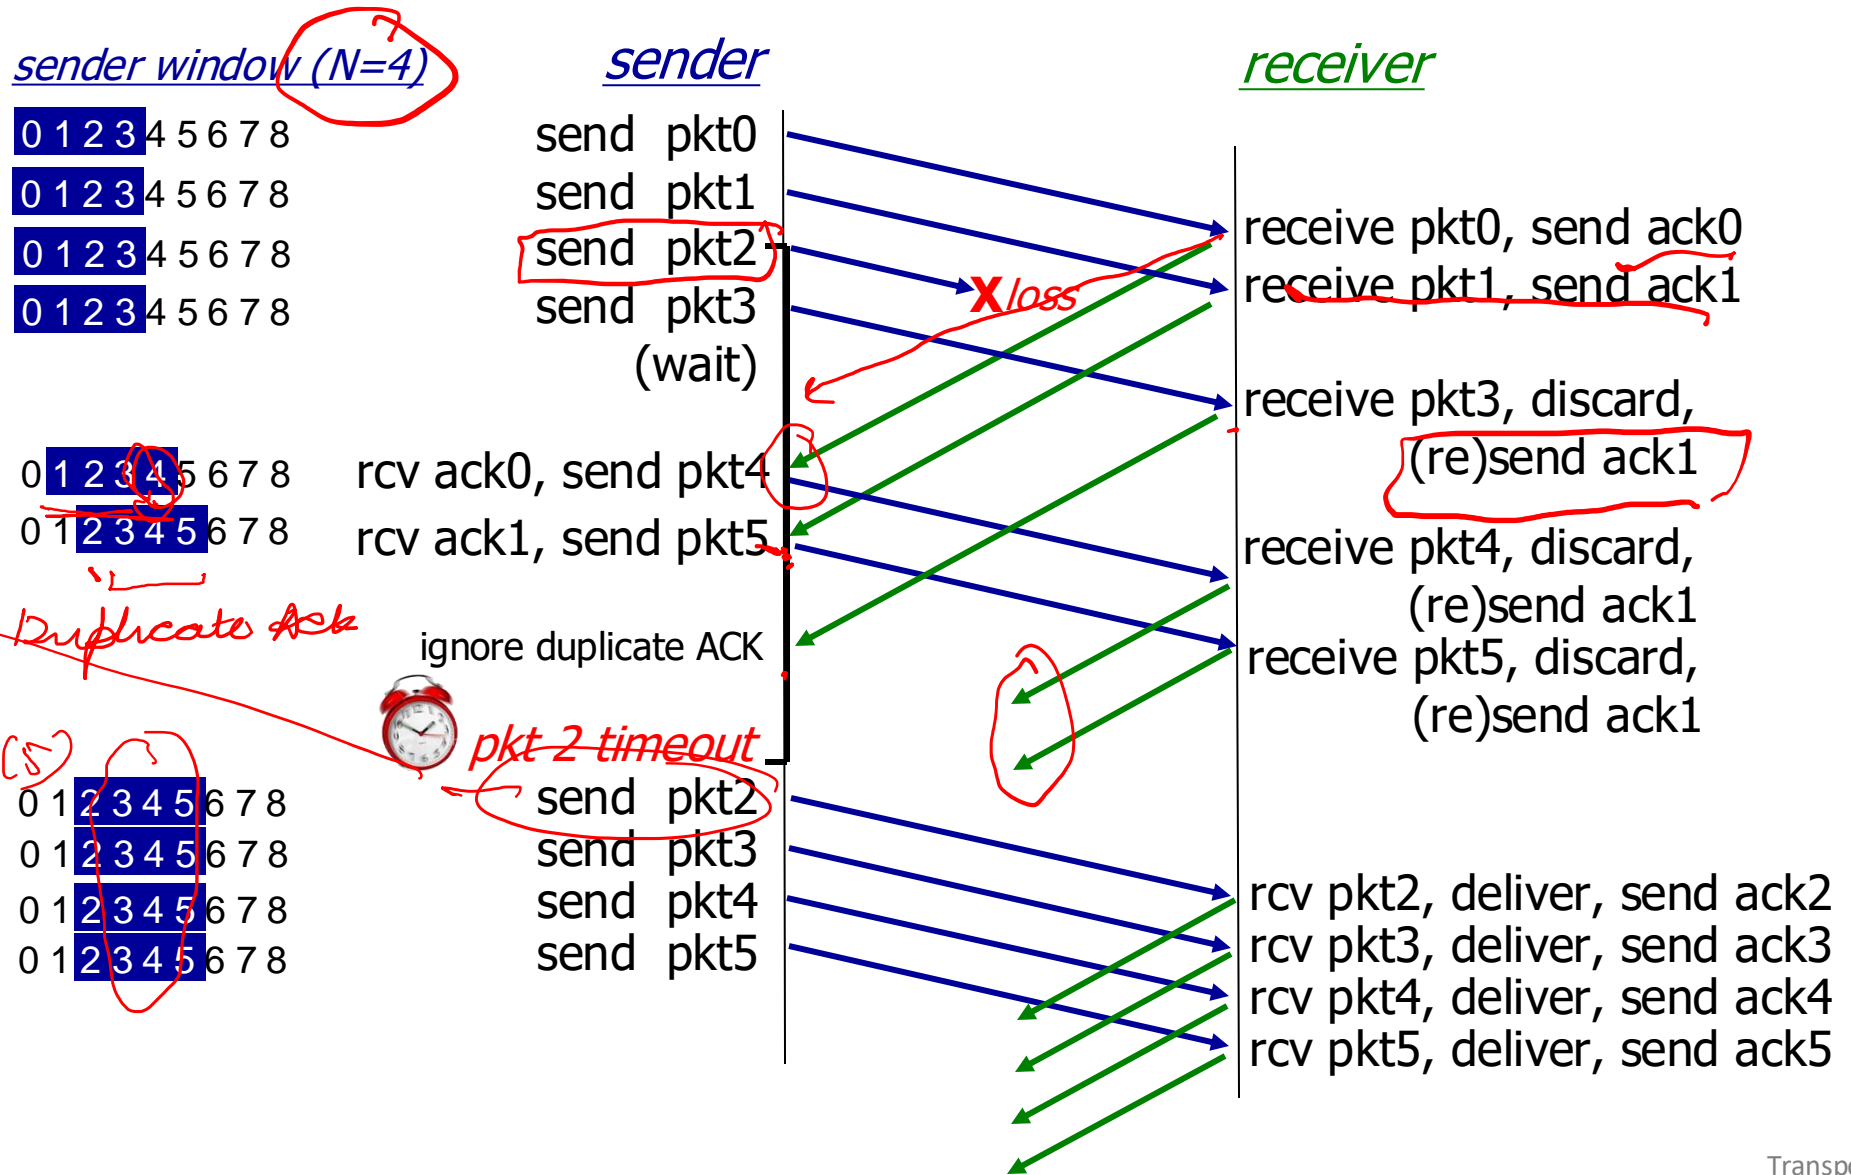
Go-Back-N: sender

- sender: “window” of up to N , consecutive transmitted but unACKed pkts
 - k -bit seq # in pkt header



- ***cumulative ACK***: $\text{ACK}(n)$: ACKs all packets up to, including seq # n
 - on receiving $\text{ACK}(n)$: move window forward to begin at $n+1$
- timer for oldest in-flight packet
- $\text{timeout}(n)$: retransmit packet n and all higher seq # packets in window


Go-Back-N in action



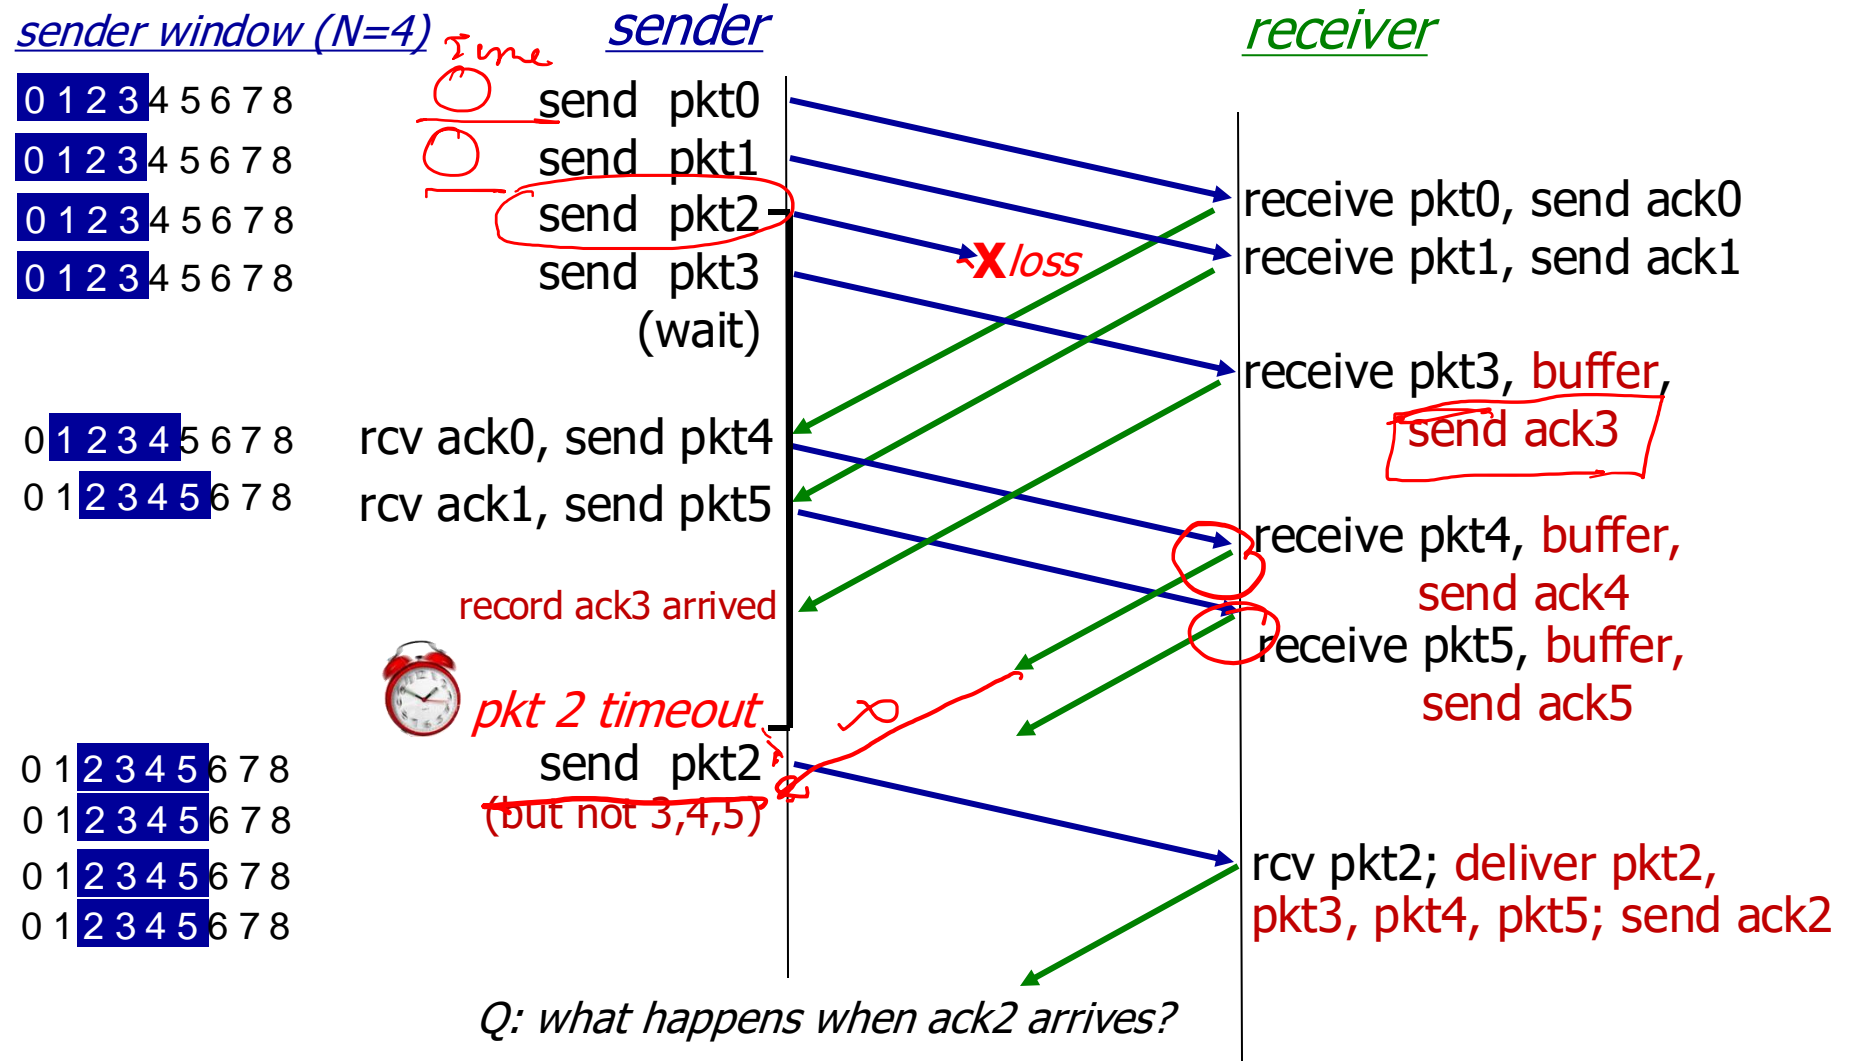
Selective repeat: the approach

- *pipelining*: multiple packets in flight
- *receiver individually ACKs* all correctly received packets
 - buffers packets, as needed, for in-order delivery to upper layer
- sender:
 - maintains (conceptually) a timer for each unACKed pkt
 - timeout: retransmits single unACKed packet associated with timeout
 - maintains (conceptually) “window” over *N* consecutive seq #s
 - limits pipelined, “in flight” packets to be within this window

Selective repeat: sender, receiver windows



Selective Repeat in action



Seq # K-bit header
 \geq sender v

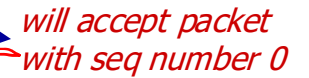
2

- sender window
(after receipt)

receiver window
(after receipt)



*will accept packet
with seq number 0*



(b) oops!

Selective repeat: a dilemma!


example:

- seq #s: 0, 1, 2, 3 (base 4 counting)
- window size=3


Q: what relationship is needed between sequence # size and window size to avoid problem in scenario (b)?

sender window
(after receipt)

receiver window
(after receipt)



- receiver can't see sender side
- receiver behavior identical in both cases!
- something's (very) wrong!



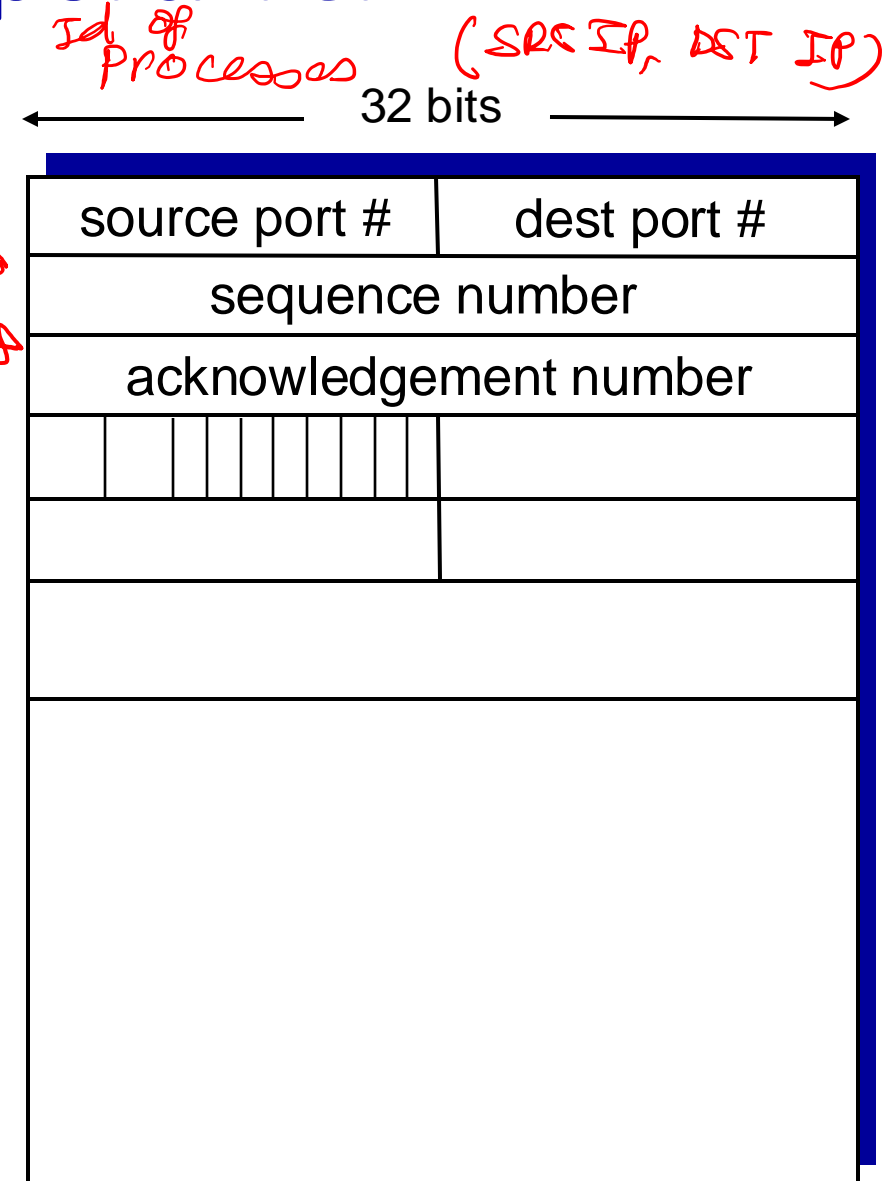

(b) oops!

Transport Control Protocol


- Connection-oriented
- Reliability and in-order delivery
- Flow control
- Congestion control

Connection-oriented transport: TCP

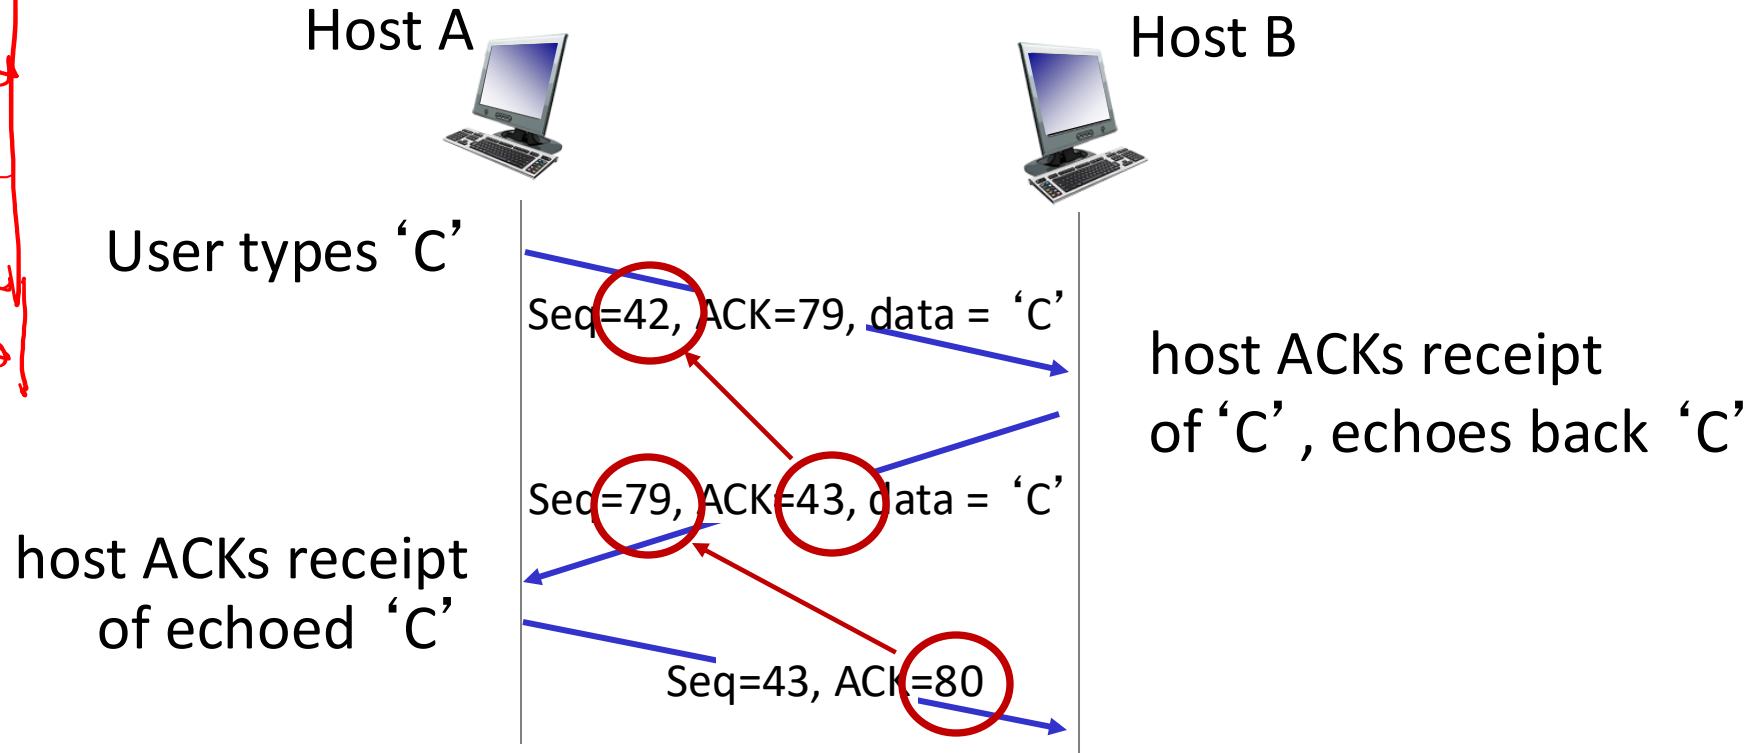
Bi-directional



TCP sequence numbers, ACKs



Randomized
Seq #



simple telnet scenario

TCP sequence numbers, ACKs

Sequence numbers:

- byte stream “number” of first byte in segment’s data

Acknowledgements:


- seq # of next byte expected from other side
- cumulative ACK

Q: how receiver handles out-of-order segments

- A: TCP spec doesn’t say, - up to implementor

outgoing segment from sender

| | |
|------------------------|-------------|
| source port # | dest port # |
| sequence number | |
| acknowledgement number | |
| | rwnd |
| checksum | urg pointer |




outgoing segment from receiver

| | |
|------------------------|-------------|
| source port # | dest port # |
| sequence number | |
| acknowledgement number | |
| | rwnd |
| checksum | urg pointer |


TCP connection management

before exchanging data, sender/receiver “handshake”:

- agree to establish connection (each knowing the other willing to establish connection)
- agree on connection parameters (e.g., starting seq #s)





```
Socket clientSocket =  
    newSocket("hostname", "port number");
```



```
Socket connectionSocket =  
    welcomeSocket.accept();
```

Agreeing to establish a connection

2-way handshake:



- Q: will 2-way handshake always work in network?
- variable delays
 - retransmitted messages (e.g. req_conn(x)) due to message loss
 - message reordering
 - can't "see" other side
- Handwritten notes in red ink include: "Req(conn)" and "Ack(conn)" with arrows, and "Req(conn)" and "Ack" with arrows, suggesting a sequence of events or a sequence of messages.

2-way handshake scenarios

TCP 3-way handshake

Client state

```
clientSocket = socket(AF_INET, SOCK_STREAM)
```

LISTEN

```
clientSocket.connect((serverName, serverPort))
```

SYNSENT

ESTAB

choose init seq num, x
send TCP SYN msg

received SYNACK(x)
indicates server is live;
send ACK for SYNACK;
this segment may contain
client-to-server data



SYNbit=1, Seq=x

SYNbit=1, Seq=y
ACKbit=1; ACKnum=x+1

ACKbit=1, ACKnum=y+1

Server state

```
serverSocket = socket(AF_INET, SOCK_STREAM)  
serverSocket.bind(('', serverPort))  
serverSocket.listen(1)  
connectionSocket, addr = serverSocket.accept()
```

LISTEN

SYN RCVD

ESTAB

choose init seq num, y
send TCP SYNACK
msg, acking SYN

received ACK(y)
indicates client is live



Closing a TCP connection

- client, server each close their side of connection
 - send TCP segment with FIN bit = 1
- respond to received FIN with ACK
 - on receiving FIN, ACK can be combined with own FIN
- simultaneous FIN exchanges can be handled

