

Computer Networks

COL 334/672

Congestion Control

Tarun Mangla

Slides adapted from KR

Sem 1, 2024-25


Quiz: Moodle

Password: tcp

Recap

- Transport Congestion Protocol
 - Connection establishment
 - Reliability
 - Flow control
 - **Congestion control**

Congestion Control



Congestion:

- informally: “too many sources sending too much data too fast for *network* to handle”
- manifestations:
 - long delays (queueing in router buffers)
 - packet loss (buffer overflow at routers)
- different from flow control!
- a top-10 problem!



congestion control:

too many senders,
sending too fast




flow control: one sender
too fast for one receiver

The Case of Congestion Collapse

could send at Round

- Early TCP protocol in 1980s used fixed size window
 - The focus was mainly on providing reliability
- As network load grew, **goodput** reduced



“costs” of congestion:

- more work (retransmission) for given receiver throughput
- unneeded retransmissions: link carries multiple copies of a packet
 - decreasing maximum achievable goodput

Need mechanisms for congestion control


Approaches towards congestion control

End-end congestion control:

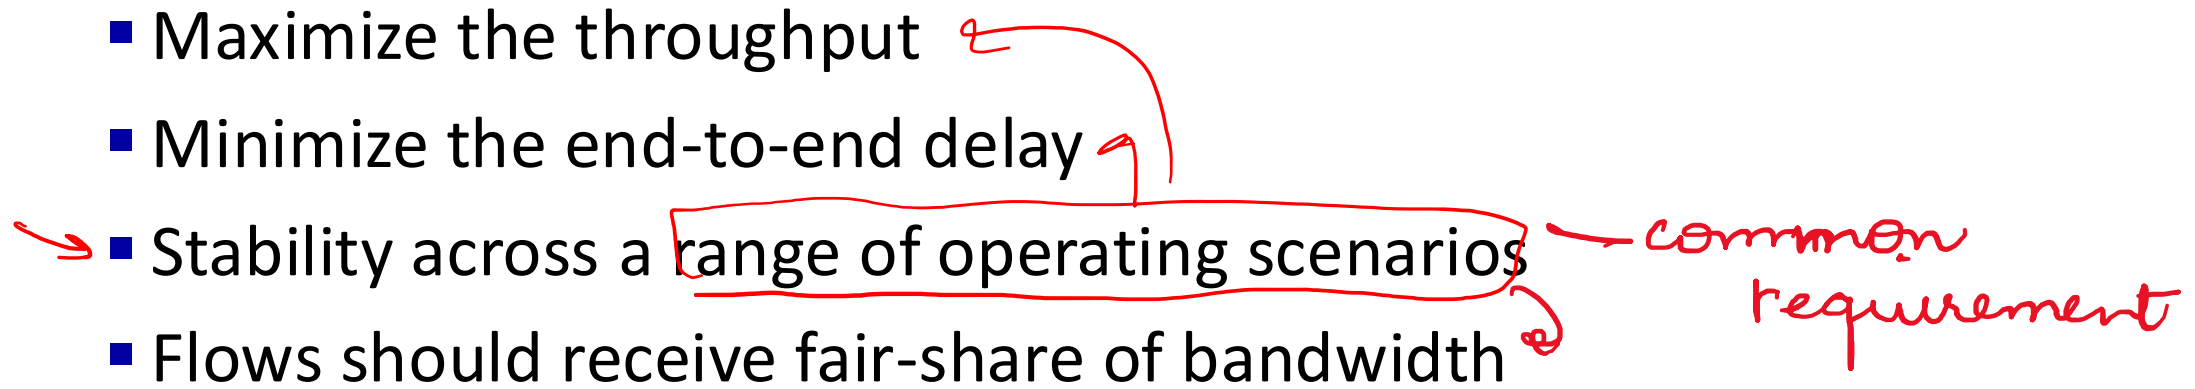
- no explicit feedback from network
- congestion *inferred*
- **approach taken by TCP**

Network-assisted congestion control:

- routers provide *direct* feedback to sending/receiving hosts with flows passing through congested router
- may indicate congestion level or explicitly set sending rate
- E.g., TCP ECN, ATM, DECbit protocols




Metrics for evaluation of a Congestion Control Algorithm?

- Maximize the throughput
 - Minimize the end-to-end delay
 - ■ Stability across a range of operating scenarios
 - Flows should receive fair-share of bandwidth
- common requirement*
- 

End-to-end congestion control

curnd

H_1



→ rate-based: send at some average rate R .
 → window-based: $(cwnd, rwnd)$
 ↓
 used by TCP

Fundamental questions

- How to determine the available ~~capacity~~ ^{bandwidth} for a flow at any time? ^{↑ cwnd}
- How to detect congestion? ^{packet loss / increase in RTT values}
- How to react to congestion once detected? ^{↓ cwnd}

TCP congestion control:

- *approach*: senders can increase sending rate until packet loss (congestion) occurs, then decrease sending rate on loss event
- Many functions possible
- *Class of linear* control algorithms:
 - $w_{i+1} = aw_i + b$
 - a and b are different in case of increase and decrease
- Additive increase, multiplicative decrease (AIMD)
 - $w_{i+1} = w_i + b$
 - $w_{i+1} = aw_i$

TCP congestion control: AIMD

- Additive increase, multiplicative decrease (AIMD)
 - Increase: $w_{i+1} = w_i + b$, Decrease: $w_{i+1} = aw_i$


Additive Increase

increase sending rate by 1
maximum segment size every
RTT until loss detected

Multiplicative Decrease

cut sending rate in half at
each loss event

AIMD sawtooth
behavior: *probing*
for bandwidth



**Why was
AIMD chosen?**

TCP fairness

Fairness goal: if K TCP sessions share same bottleneck link of bandwidth R , each should have average rate of R/K

