

Computer Networks

COL 334/672

Congestion Control


Tarun Mangla

Slides adapted from KR

Sem 1, 2024-25

Recap: TCP Congestion Control

- **What:** Adjust sending rate at end-host to avoid congestion in the network
- **Why:** congestion → wasted network resources
- **How:**
 - **Design Space:** Multiple approaches possible. TCP chose:
 - End-to-end
 - Distributed
 - Window-based
 - **General approach:**
 1. Probe b/w by increasing window
 2. Detect congestion
 3. Backoff by decreasing window
 - **Multiple ways:** AIMD is chosen (desirable stability properties, provides fairness)



Window Update in AIMD

$0 < \beta < 1$


AIMD: AI: $wnd = wnd + \alpha$
 $wnd \max \rightarrow \alpha > 1 \Rightarrow -1$ MD: $wnd = \frac{wnd}{\beta}$

- **Goal:** For additive increase, update window size by 1 every RTT. But when?

$$wnd = wnd + \frac{1}{wnd}$$


goal 1MSS \uparrow per RTT

$$wnd = wnd + \left(\frac{MSS}{wnd} \right) \times MSS$$




Window Update in AIMD

- **Goal:** For additive increase, update window size by 1 every RTT. But when?
 - Window updated as ACKs arrive
 - How much should be the update?
-
- Because TCP updates based on ACKs, TCP is self-clocking **Benefit: Self-clocking smooths packet sending rate**



TCP Self Clocking

- ACK clocking smooths packet sending rate






AIMD is TCP Fair

Active Queue Management (AQM)
↳ Random Early Drop (RED)

Example: two competing TCP sessions: ②. Is this mechanism TCP friendly scenarios?

AIMD: is not RTT - fair



Beginning of the Connection

$cwnd = 10^4$ $IN = 1$
 $RTT = 10^{-2} s$
How much time
= 10^2

- What *cwnd* size should we start with?

- **Goal:** We want to quickly near the *right* rate.

- A linear increase with a small value of *cwnd* is painfully slow!

↓
exponential

↳ start w/ a large value?


- What are the options?

- Start with a large value of *cwnd*
- Start with a small value of *cwnd* but increase it faster

TCP slow start

Start with a large IW : Rwnd

- when connection begins, increase rate exponentially until first loss event:
 - initially **cwnd** = 1 MSS
 - double **cwnd** every RTT
 - done by incrementing **cwnd** for every ACK received
- *summary*: initial rate is slow, but ramps up exponentially fast

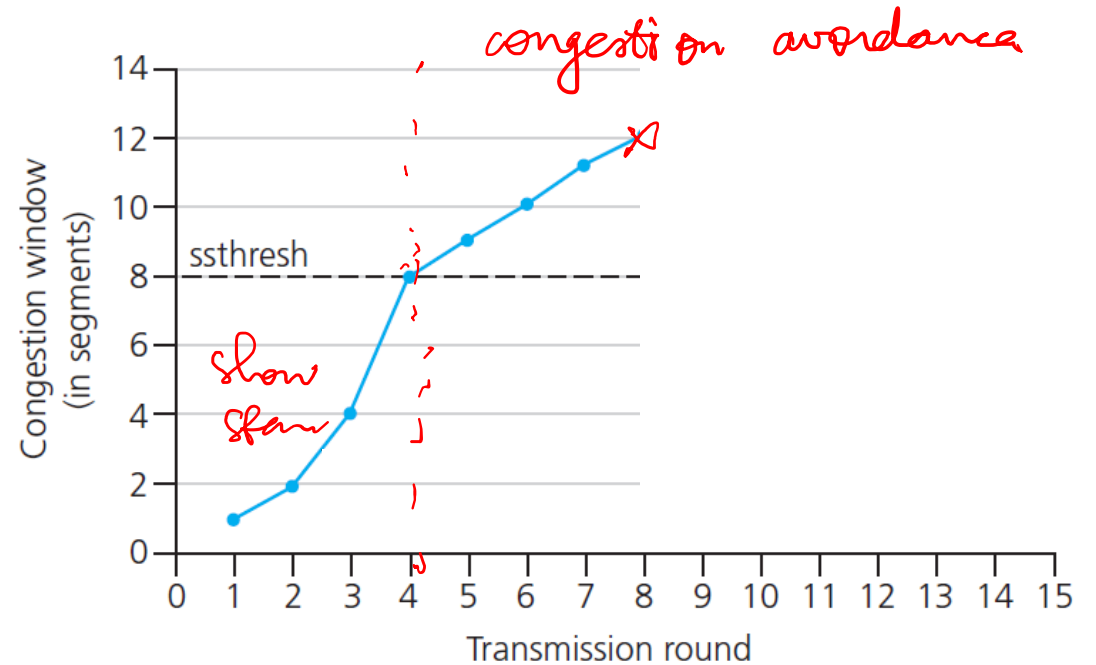


Until when?

TCP Slow Start and Congestion Avoidance

- Uses a threshold, ssthresh, after which TCP enters congestion avoidance
- *What happens when a loss occurs?*

$IN \geq 1 \rightarrow 2$



What happens when a loss occurs?

Timeout
Triple dup
Fast Retransmit
Ack


- When loss is due to **triple duplicate ACK**: congestion is not severe!

→ $ssthresh = 6$
 $cwnd = 6$

- Set $ssthresh$ to $cwnd/2$
- On receipt of another duplicate ACK, send 1 new segment
- Once a new ACK arrives, set $cwnd = ssthresh$ (or $cwnd/2$)
- Enter **congestion avoidance** phase

TCP Reno: Fast retransmit, fast recovery!

TCP Tahoe:




What happens when a loss occurs?

- When loss is due to timeout:

severe congestion!! *ssthresh = 6*

- Set ssthresh to cwnd/2
- Reset cwnd to 1, *IW*
- Enter slow start phase




What happens when a loss occurs?

- Loss is due to timeout: **severe congestion!!**
 - Set *ssthresh* to $cwnd/2$
 - Reset *cwnd* to 1
 - Enter **slow start phase**
- Loss is due to triple duplicate ACK: **congestion is not severe!**
 - Set *ssthresh* to $cwnd/2$
 - On receipt of another duplicate ACK, send 1 new segment
 - Once a new ACK arrives, set *cwnd* = *ssthresh* (or $cwnd/2$)
 - Enter **congestion avoidance** phase

TCP Reno Throughput: Macroscopic Description

- Throughput: area under the curve

$$\text{Average Tput} : \frac{3W}{4}$$




Inefficient for networks with high bandwidth delay product!

Can we do faster?

Is there a better way than AIMD to “probe” for usable bandwidth?

- Insight/intuition:

- W_{\max} : sending rate at which congestion loss was detected
- congestion state of bottleneck link probably (?) hasn't changed much
- after cutting rate/window in half on loss, initially ramp to to W_{\max} *faster*, but then approach W_{\max} more *slowly*



TCP CUBIC

- K: point in time when TCP window size will reach W_{\max}
 - K itself is tunable
- increase W as a function of the *cube* of the distance between current time and K

$$W(t) = C(t - K)^3 + W_{\max}$$

$$K = \sqrt[3]{\frac{W_{\max}\beta}{C}}$$

- TCP CUBIC default
in Linux, most
popular TCP for
popular Web
servers

Attendance

