

Computer Networks

COL 334/672

Congestion Control

Tarun Mangla

Slides adapted from KR

Sem 1, 2024-25

Recap: TCP

- Connection establishment
- Reliability
- Flow control
- Congestion control
 - Loss-based
 - Delay-based
 - Network-assisted

Evolving transport-layer functionality

- TCP, UDP: principal transport protocols for 40 years
- different “flavors” of TCP developed, for specific scenarios:

>4 papers out of 62 papers
on TCP in Sigcomm24

Scenario	Challenges
Long, fat pipes (large data transfers)	Many packets “in flight”; loss shuts down pipeline
Wireless networks	Loss due to noisy wireless links, mobility; TCP treat this as congestion loss
Long-delay links	Extremely long RTTs
Data center networks	Latency sensitive
Background traffic flows	Low priority, “background” TCP flows

Sharing the network

Session Chair: Prateesh Goyal (Microsoft Research)

Keeping an Eye on Congestion Control in the Wild with Nebby [Research Track](#)

Ayush Mishra (National University of Singapore); Lakshay Rastogi (Indian Institute of Technology, Kanpur); Raj Joshi, Ben Leong (National University of Singapore)



.SUSS: Improving TCP Performance by Speeding Up Slow-Start [Research Track](#)

Mahdi Arghavani, Haibo Zhang, David Eysers (School of Computing, University of Otago, New Zealand); Abbas Arghavani (School of Innovation, Design and Engineering, Mälardalen University, Sweden)



Principles for Internet Congestion Management

[Research Track](#)

Lloyd Brown (UC Berkeley); Albert Gran Alcoz (ETH Zürich); Frank Cangialosi (BreezeML); Akshay Narayan (Brown University); Mohammad Alizadeh, Hari Balakrishnan (MIT); Eric Friedman (ICSI and UC Berkeley); Ethan Katz-Bassett (Columbia University); Arvind Krishnamurthy (University of Washington); Michael Schapira (Hebrew University of Jerusalem); Scott Shenker (ICSI AND UC Berkeley)




CCAnalyzer: An Efficient and Nearly-Passive Congestion Control Classifier [Research Track](#)

Ranysha Ware, Adithya Abraham Philip (Carnegie Mellon University); Nicholas Hungria (Carnegie Mellon University); Yash Kothari, Justine Sherry, Srinivasan Seshan (Carnegie Mellon University)

Internet Layered Architecture

- *application*: supporting network applications
 - HTTP, IMAP, SMTP, DNS
- *transport*: process-process data transfer
 - TCP, UDP
- *network*: routing of datagrams from source to destination
 - IP, routing protocols
- *link*: data transfer between neighboring network elements
 - Ethernet, 802.11 (WiFi), PPP
- *physical*: bits “on the wire”



Networked Applications

- How to write a networked application?

Using socket programming

- Multiple networked applications

Can you give examples?

N/w Apps

① File server

② Email

③ Video streaming

④ Web browsing

⑤ Torrent

⑥ ssh

⑦ whatsapp / chatting app

⑧ Video conference

Popular Applications

- HTTP → *Web browsing*
- Email
- DNS → *Infrastructure apps*
- P2P
- Video Streaming

HTTP

- Overview of HTTP
- Request/response message format
- State management
- Caching
- Request pipelining

Web and HTTP

- World Wide Web or Web was the 2nd **killer app** over the Internet
- **Goal of Web**: organize and retrieve information over Internet
- Web is based on two sister protocols: **HTML** and **HTTP**
- HTML: HyperText Markup Language
 - language to create and structure web page
 - web page consists of *objects*, each of which can be stored on different Web servers
 - **base HTML-file** which includes several referenced objects, each addressable by a URL

`www.wikipedia.org/wiki/http`

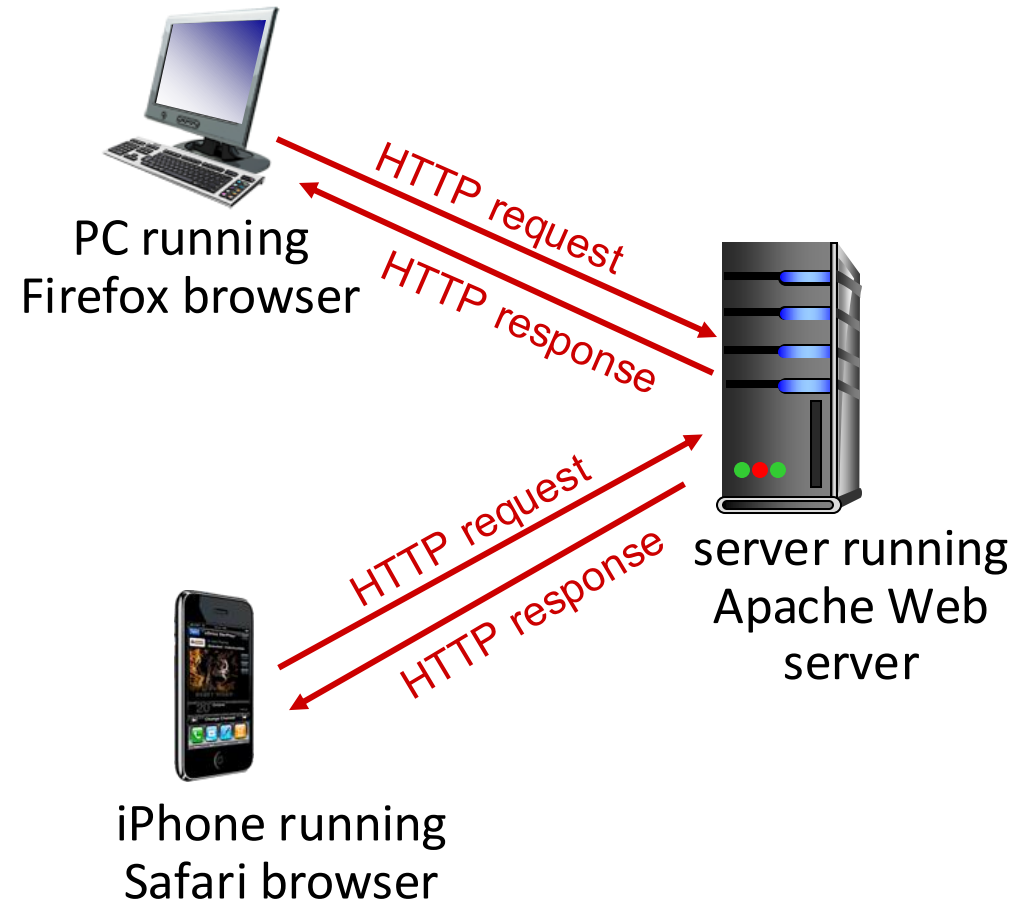
host name

path name

HTTP overview

HTTP: hypertext transfer protocol

- Web's application-layer protocol
- Request/response protocol or client/server protocol
 - *client*: browser that requests, receives, (using HTTP protocol) and “displays” Web objects
 - *server*: Web server sends (using HTTP protocol) objects in response to requests



HTTP overview (continued)

*Which transport protocol? **TCP***

Downloading a webpage

How to download multiple web objects?

Non-persistent HTTP

1. TCP connection opened
2. at most one object sent over TCP connection
3. TCP connection closed

downloading multiple
objects required multiple
connections

Non-persistent HTTP: example

User enters URL: `www.someSchool.edu/someDepartment/home.index`
(containing text, references to 10 jpeg images)



1a. HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80



1b. HTTP server at host `www.someSchool.edu` waiting for TCP connection at port 80 “accepts” connection, notifying client

2. HTTP client sends HTTP *request message* (containing URL) into TCP connection socket. Message indicates that client wants object `someDepartment/home.index`

3. HTTP server receives request message, forms *response message* containing requested object, and sends message into its socket

time



Non-persistent HTTP: example (cont.)


User enters URL: `www.someSchool.edu/someDepartment/home.index`
(containing text, references to 10 jpeg images)



5. HTTP client receives response message containing html file, displays html. Parsing html file, finds 10 referenced jpeg objects

6. Steps 1-5 repeated for each of 10 jpeg objects

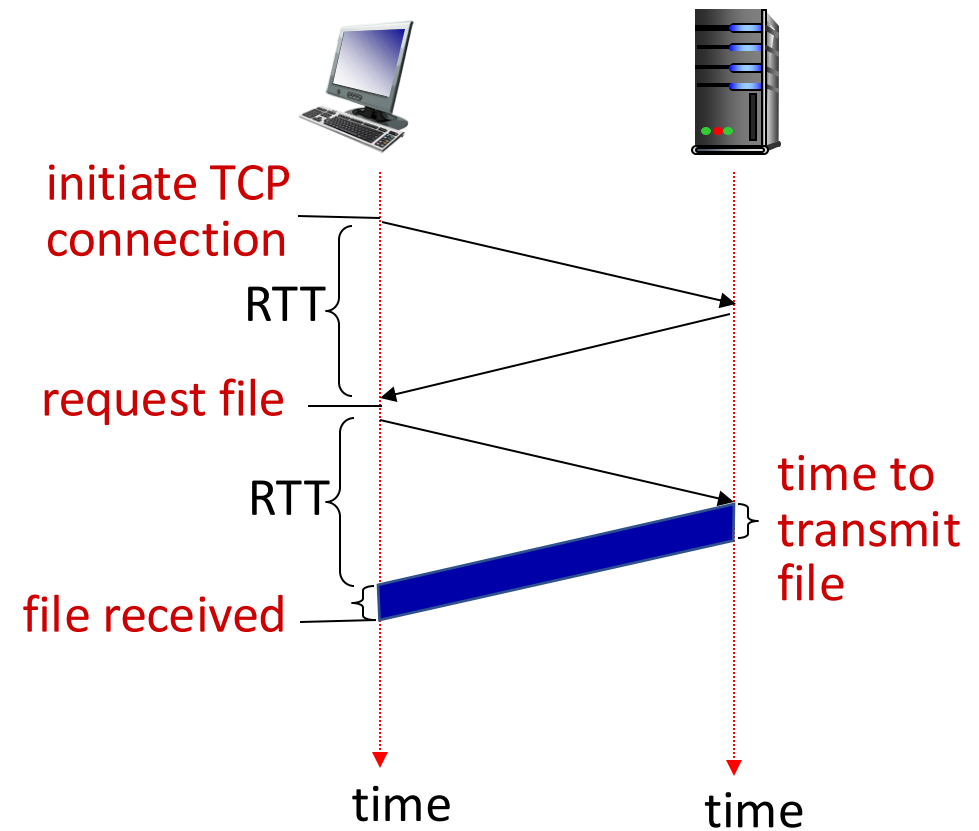
4. HTTP server closes TCP connection.



Non-persistent HTTP: response time

HTTP response time (per object):

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- object/file transmission time



Non-persistent HTTP response time = 2RTT + file transmission time

Persistent HTTP (HTTP 1.1)

Non-persistent HTTP issues:

- requires 2 RTTs per object
- OS overhead for *each* TCP connection
- browsers often open multiple parallel TCP connections to fetch referenced objects in parallel

Persistent HTTP (HTTP1.1):

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects (cutting response time in half)

HTTP

- Overview of HTTP
- **Request/response message format**
- State management
- Caching
- Request pipelining

HTTP request message


- two types of HTTP messages: *request, response*
- HTTP request message:
 - ASCII (human-readable format)

request line (GET, POST,
HEAD commands) →

/ carriage return character
/ line-feed character

carriage return, line feed →
at start of line indicates
end of header lines

HTTP request message: general format



Other HTTP request messages

POST method:

- web page often includes form input
- user input sent from client to server in entity body of HTTP POST request message

GET method (for sending data to server):

- include user data in URL field of HTTP GET request message (following a '?'):

`www.somesite.com/animalsearch?monkeys&banana`


HEAD method:

- requests headers (only) that would be returned *if* specified URL were requested with an HTTP GET method.

PUT method:

- uploads new file (object) to server
- completely replaces file that exists at specified URL with content in entity body of POST HTTP request message

HTTP response message

status line (protocol  HTTP/1.1 200 OK
status code status phrase)

HTTP response status codes

- status code appears in 1st line in server-to-client response message.
- some sample codes:

200 OK

- request succeeded, requested object later in this message

301 Moved Permanently

- requested object moved, new location specified later in this message (in Location: field)

400 Bad Request

- request msg not understood by server

404 Not Found

- requested document not found on this server

505 HTTP Version Not Supported

Attendance

